

# Boxes in Boxes

## Virtualisation and Containerisation in the Context of Embedded routers

Wouter Cloetens & Wojtek Makowski

OpenWrt Summit  
Berlin, 2016-10-13

rev. 2

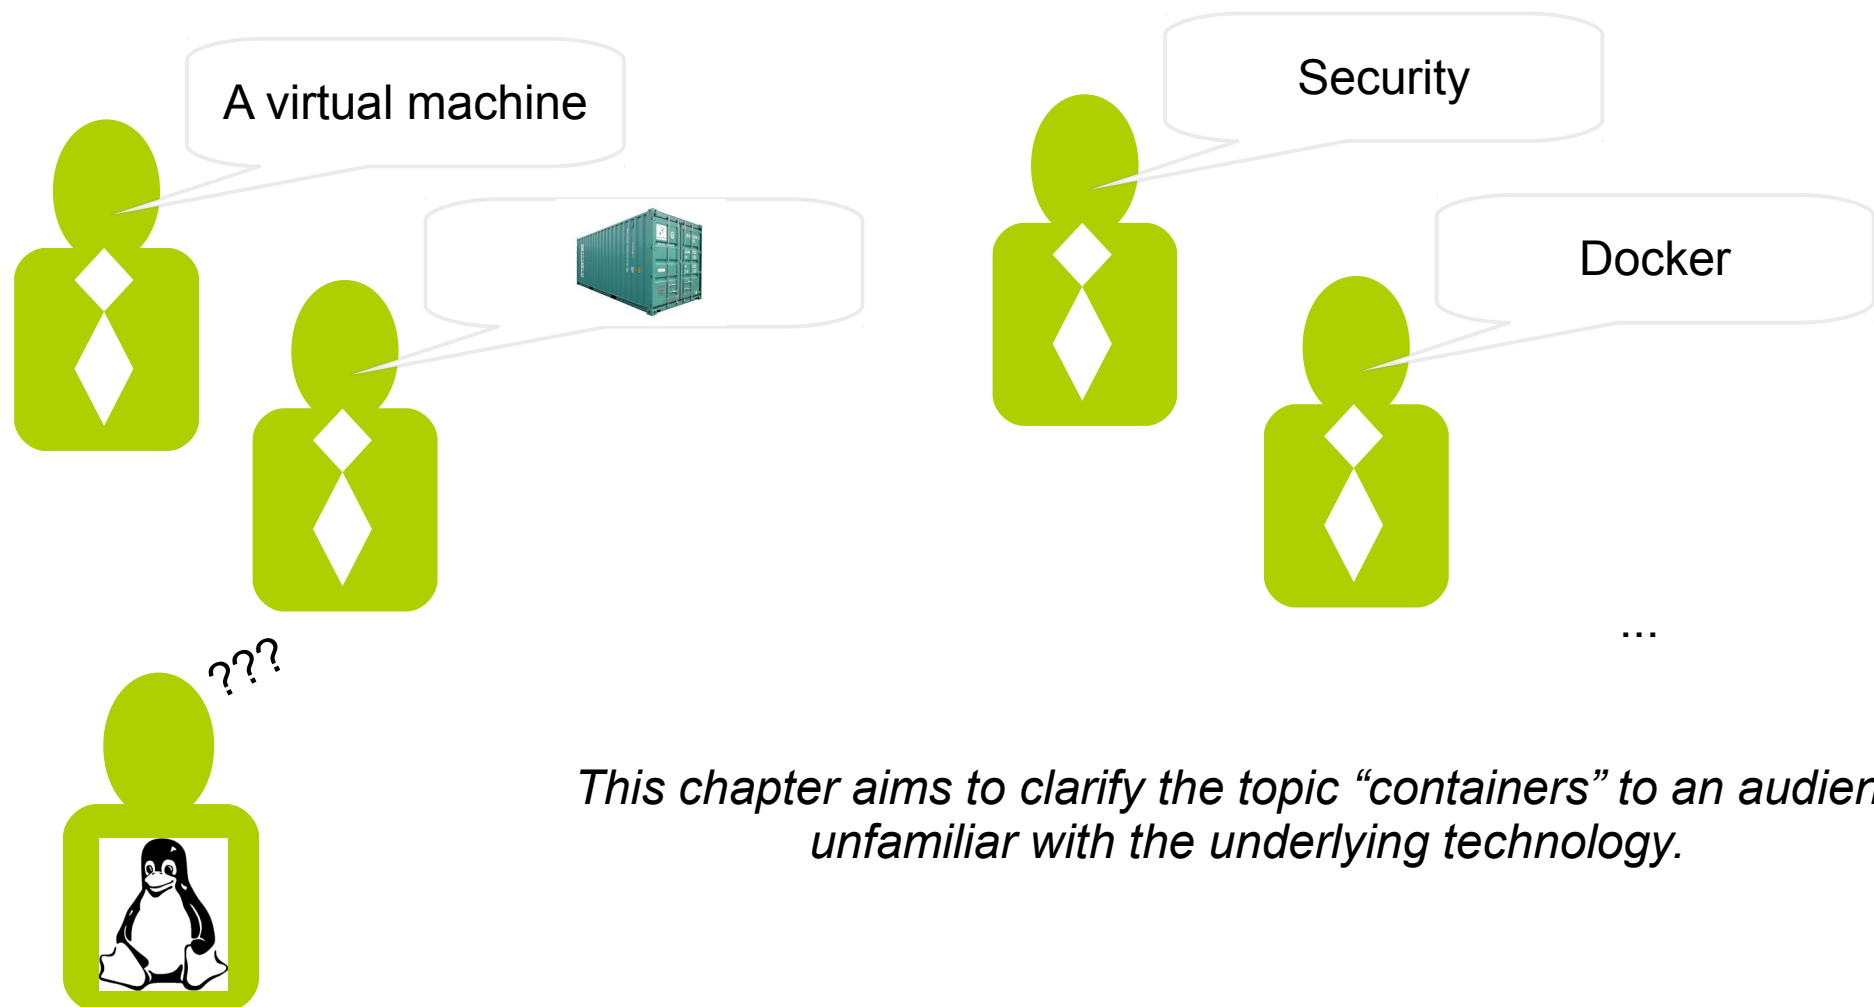


# **Part 1**

## **Introduction to Containers**

Everyone talks about containers...

... but what exactly is a Linux 'container' ?



*This chapter aims to clarify the topic “containers” to an audience unfamiliar with the underlying technology.*

## Common misconceptions

- Linux containers are virtual machines

*No, they are not.*

- Linux containers can run on any platform

*No, they can only run on a platform running the Linux kernel*

- Running something in a container makes it completely secure


*No*

- Containers will solve all my (security) problems

*No, although they are really powerful, they have limitations*

# Prerequisites for running Linux containers ?

- A (recent) Linux kernel with various security features enabled  
[preferably > Linux kernel 4.1]
- A framework to run containers : Docker, LXC, ...



## Container frameworks:

There are various technical frameworks that can be used for creating, managing and running « containers »:


*Docker, LXC, ...*

Each of these frameworks has its own features, but they all use the same underlying « security » mechanisms implemented in the Linux kernel :

Namespaces, Cgroups, Seccomp, Capabilities, user/group, chroot, apparmor, SELinux

*⇒ This presentation will describe the common security mechanisms used by containers and illustrate this with examples.*

## Linux user space process running as root:



- Has access to complete file system
- Has access to all device nodes
- Can perform any kernel call
- Can reconfigure the network
- Can potentially consume all resources
- Can send signals to other processes (SIGKILL)
- ...

## Linux user space process running as root:



root 1107 /bin/sh


*⇒ If something goes wrong either due to a bug or due to a hacking attempt, nothing shields the rest of the system!*

```
rm -rf $(MY_UNDEFINED_VAR)/*
```

```
while (true) {  
  switch(case){  
    Up: break  
    Default: continue  
  }  
}
```




## Step 1: Run the process as unprivileged user



- Discretionary Access Control: standard Linux (file) permissions and limitations are applied
- More limited access and control of devices
- More limited access and control of network devices
- Permissions of all kernel calls are checked prior to executing them
- Limited ability to send signals to processes owned by different users (e.g. !SIGKILL)

## Can it still be made more secure?



**YES!**

⇒ *A process might still have potentially 'dangerous' capabilities!*

→ *set UID's on files*


→ *Open RAW sockets*

→ *Perform network configuration*

→ ...

setsockopt() is OK  
...  
Using raw socket and UDP protocol  
...


## Step 2: Limit capabilities of the process



- Defined per thread
- Examples of available capabilities :

CAP_CHOWN	CAP_NET_ADMIN	CAP_SYS_ADMIN
CAP_FOWNER	CAP_NET_RAW	CAP_SYS_BOOT
CAP_FSETID	CAP_IPC_LOCK	CAP_SYS_NICE
CAP_KILL	CAP_IPC_OWNER	CAP_SYS_RESOURCE
CAP_SETGID	CAP_SYS_MODULE	CAP_SYS_TIME
CAP_SETUID	CAP_SYS_RAWIO	CAP_MKNOD
CAP_MAC_ADMIN	CAP_SYS_CHROOT	CAP_MAC_OVERRIDE
- Fewer capabilities = Better security

# Can it still be made more secure?



**YES !**


⇒ *The process can still see all processes, network interfaces, mount points, IPC sockets, cgroups, users and UTS configuration!*

```
>ps
PID USER      COMMAND
0   root       /sbin/init
```

```
>ifconfig
lo      Link encap:Local Loopback
        inet addr:127.0.0.1  Mask:255.0.0.0
        inet6 addr: ::1/128 Scope:Host
        UP LOOPBACK RUNNING  MTU:65536  Metric:1
```

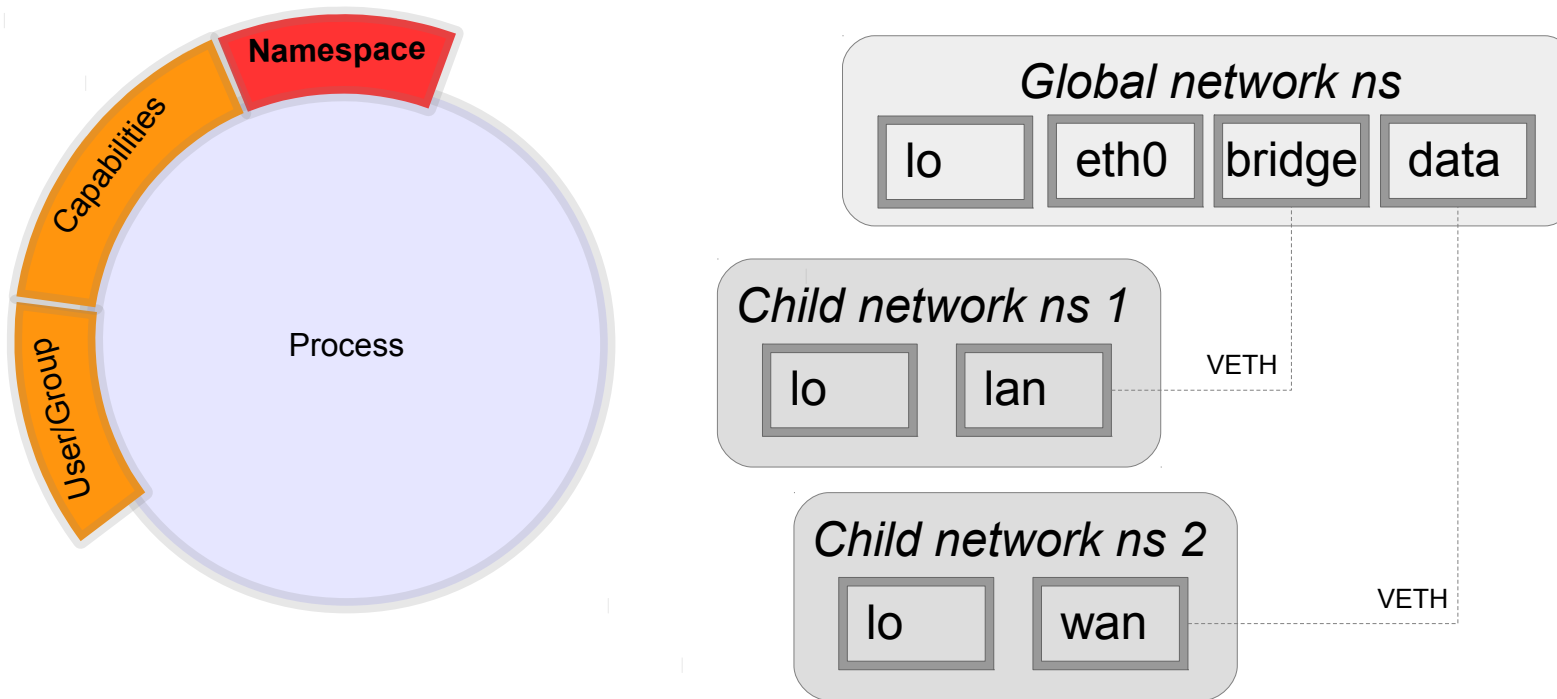
```
socket.error: [Errno 98] Address already in use
```

## Step 3: Apply namespaces



- The kernel currently supports 6 different namespaces :
  - Network
  - PID/Process
  - Mount
  - User
  - IPC
  - UTS
- Provides isolation for each namespace type
- More namespace types are under development

## Step 3: Apply namespaces: Network namespace



*In global network namespace:*

```
> ifconfig
lo      Link encap:Local Loopback
eth0    Link encap:Ethernet HWaddr 20:47:47:73:e1:52
bridge  Link encap:Ethernet HWaddr 10:38:1e:1a:ea:33
data    Link encap:Ethernet HWaddr 22:50:e5:a0:f0:65
```


*In child network namespace 1:*

```
> ifconfig
lo      Link encap:Local Loopback
lan     Link encap:Ethernet HWaddr 56:87:47:73:45:11
```

*In child network namespace 2:*

```
> ifconfig
lo      Link encap:Local Loopback
wan     Link encap:Ethernet HWaddr 02:14:a5:e4:4f:15
```

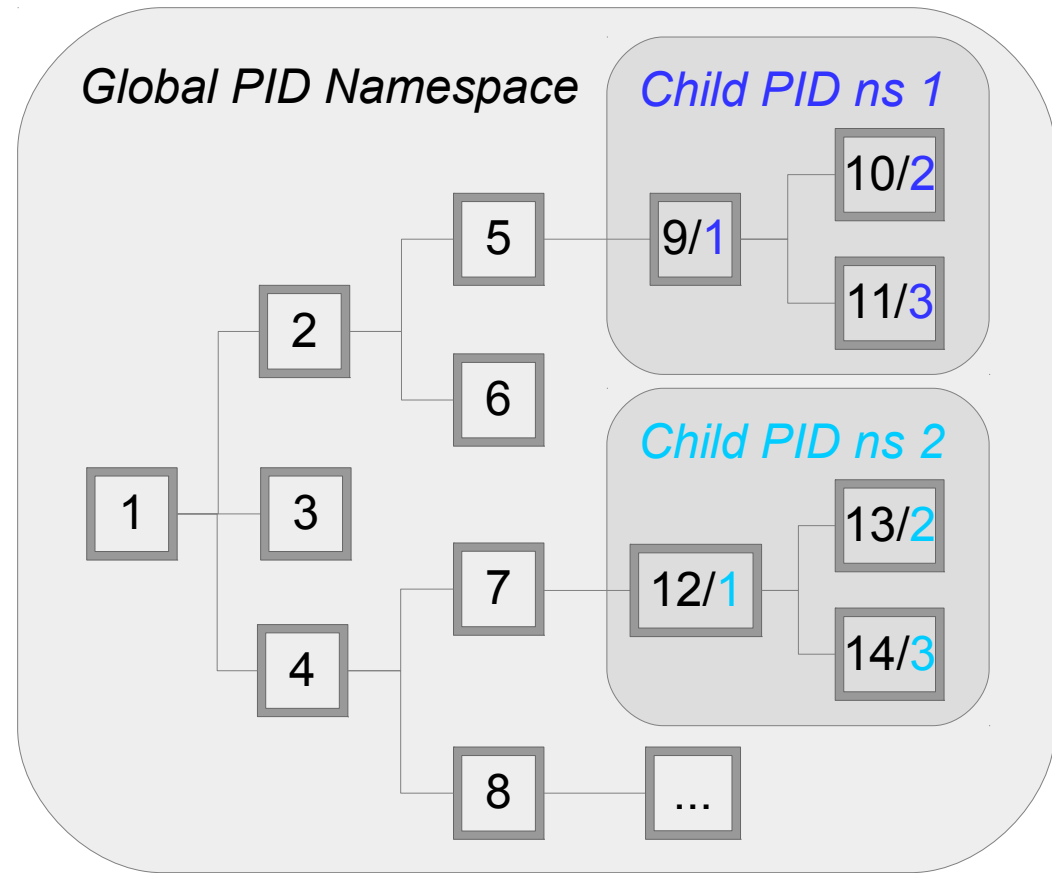
## Step 3: Apply namespaces: Process namespace



*In global mount namespace:*

> ps -A

PID	CMD	PID	CMD
1	init	8	rcu_sched
2	kthreadd	9	httpd
3	rcuos/0	10	cgi-bin
4	softirqd/0	11	php
5	kworker/0:0H	12	node.js
6	kworker/0:1H	13	/bin/sh
7	rcuob/0	14	mongod



*In PID ns 1:*

$$> ps -A$$

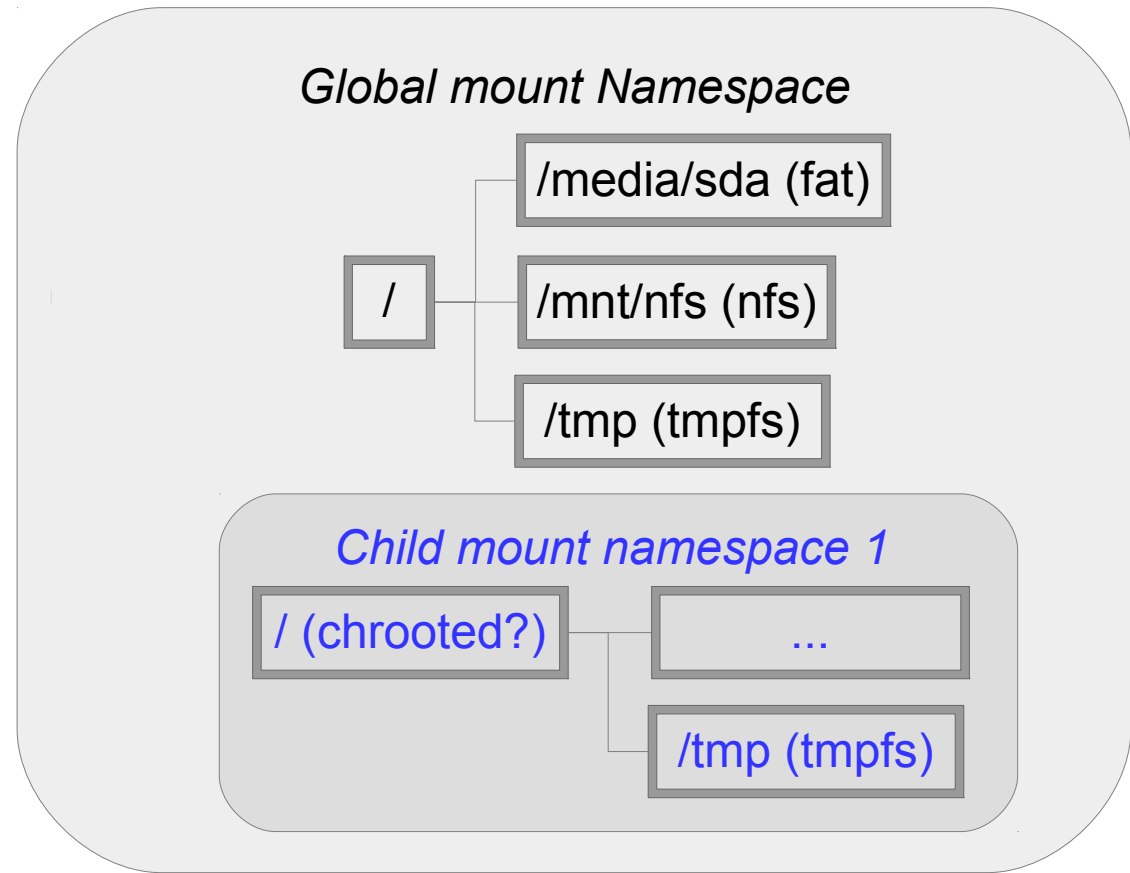

PID	CMD
1	httpd
2	cgi-bin
3	php

*In PID ns 2:*

$$> ps -A$$

PID	CMD
1	node.js
2	/bin/sh
3	mongodb

## Step 3: Apply namespaces: Mount namespace



*In global mount namespace:*


```
> mount
/dev/sda3 on / type ext4 (rw,errors)
10.0.0.2:/home on /mnt/nfs type nfs(rw)
none on /tmp type tmpfs(rw)
```

*In mount ns 1:*

```
> mount
none on / type ext4 (rw,errors)
none on /tmp type tmpfs(rw)
```



## Step 3: Apply namespaces: User namespace



- User ID's (UID) and group ID's (GID) are specific per namespace
- Each UID/GID in another namespace maps onto a real UID/GID in the host
- Process can be privileged user within its own namespace (e.g. run with root-like privileges, without impacting the host)


*In global user namespace:*

```
>cat /etc/passwd  
root:x:0:0:root:/root:/bin/bash  
user1:x:1:1:user1:/home/user1:/bin/bash  
user2:x:2:2:user2:/home/user2:/bin/bash  
virtuser1:x:3:3:virtuser1:/home/virtuser1:/bin/bash  
virtuser2:x:4:4:virtuser2:/home/virtuser2:/bin/bash
```

*In different user ns 1:*


```
> cat /etc/passwd  
root:x:0:0:root:/root:/bin/bash  
service:x:1:1:service:/service:/bin/bash
```

## Step 3: Apply namespaces: IPC namespace




- Provides namespace separation for the following IPC mechanisms :
  - System V IPC (svipc)
  - POSIX message queues (mq\_overview)
- Processes in different IPC namespace will not be able to communicate using the above IPC mechanisms.

## Step 3: Apply namespaces: UTS namespace



- Provides namespace separation for some system identifiers:
  - nodename
  - domainname
- Processes in different UTS namespaces will be able to use the identifiers without impacting the rest of the system.

# Can it still be made more secure?




**YES!**

⇒ *How can you protect against abuse of system resources? (e.g. consuming all CPU/RAM)*

PID	USER	PR	NI	VIRT	RES	SHR S	%CPU	%MEM	TIME+ COMMAND
2064	user1	20	0	2765836	1.029g	115464 S	<b>97.0</b>	6.6	75:25.45 firefox
1337	root	20	0	488716	15.029g	31828 S	3.0	<b>93.4</b>	7:51.26 Xorg

## Step 4: Apply CGroups




```
cgroup.memory.limit_in_bytes = 512M  
(apply on a process that consumes all memory)
```

*Out of memory:  
Kill process 29957 (firefox)  
score 366 or sacrifice child*

- CGroup features :
  - Resource limiting – groups can be set to not exceed a configured memory limit, which also includes the file system cache
  - Prioritisation – some groups may get a larger share of CPU utilisation or disk I/O throughput
  - Accounting – measures a group's resource usage, which may be used, for example, for billing purposes
  - Control – freezing groups of processes, their checkpointing and restarting
- CGroup restrictions can be applied to both individual processes and groups of processes

## Can it still be made more secure?



**YES!**

⇒ *The process still has access to the complete file system :*


→ *One misconfigured permission...*

→ *It can still launch most binaries*

Drwxrwxrwx user group /etc/config


/bin/curl  
/bin/sh  
/bin/...

## Step 5: Chroot: Limit access to the file system



- Use chroot to only allow access to files that are really needed (libraries, executables, configuration files, data files)
- From a process point of view: only a very small file system is visible.
- Only root can escape a chroot environment
- Different possibilities to implement the alternative file system: hard links, overlays, etc.
- Considered by most people as a 'hardening' feature, not a 'security' feature as you can still potentially escape the chrooted file system.

## Can it still be made more secure?



**YES!**

Chroot shields the file system, but is this enough ?

→ one could still escape a chroot


→ difficult to maintain

⇒ The kernel supports several alternative ways to configure fine-grained access control per process, using Mandatory Access Control:

- SELinux
- AppArmor
- Smack
- Tomoyo




## Step 6A: SELinux



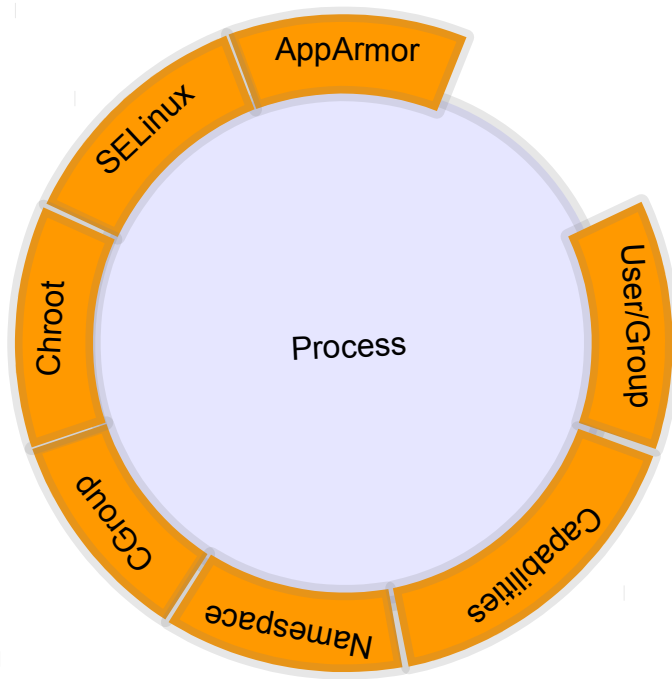
- Access control per process
- Extends traditional Linux security mechanisms
- SE Rules allows (ultra) fine grained control
- Applies labels to files
- Perceived by many people as complex to learn and maintain
- Can be used in combination with chroot

## Step 6B: AppArmor



- Alternative to SELinux
- Access control per process
- Extends traditional Linux security mechanisms
- Per-program profile
- Profiles allows fine grained control per file path
- Automatic learning mode
- Can be used in combination with chroot


## Can it still be made more secure?



**YES!**


⇒ Depending on the application you want to run, you might want to consider putting it into a 'sandbox'.

## Step 7: Seccomp (secure computing mode)




- Not applicable to all processes
- One way transition to 'secure' mode
- Allowed system calls :
  - exit
  - sigreturn
  - read/write on open file descriptor
- Any other system call results in SIGKILL
- This imposes a 'limitation', not 'virtualisation'

## Imagine Process = /bin/sh, what do you see ?



- A kind of 'virtualised' shell:
  - Virtual users (user Namespace)
  - Only one active process (PID Namespace)
  - Virtual network devices (network Namespace)
  - Limited availability of system calls (Capabilities, Seccomp)
  - Limited set of files visible (chroot)
  - Limited set of files accessible (User/Group, SELinux/APPArmor)
  - Limited use of system resources (CGroup)

## Imagine Process = init, what do you have?



- A lightweight 'virtual machine':
  - Using a shared kernel
  - Having bare-metal speed
  - Having its own file system
  - Having limited access to system calls
  - Having limited access to devices
  - Having limited access to system resources

## What added value do container frameworks like Docker/LXC offer?


- One configuration file per process or set of processes
- User space tools to manage the configuration
- User space tools to start/stop/freeze/load 'containers'
- Utilities to (remotely) manage/deploy containers
- ...

## **Part 2**

# **Containers vs. Virtual Machines**



## Can it still be made more secure?




**YES!**

⇒ We're still sharing the same kernel...

- the same device drivers
- risk of currently unknown security holes allowing privilege escalation
- CGroups accounting may not be completely accurate
- all hardware resources *must* be supported by the same kernel

## Step 8: Hardware Virtualisation



# Containers or Hardware Virtualisation?

	Benefit	Containers	HW Virtualisation
1	Split firmware into several independent entities	**	**
2	Ensure runtime independence	*	**
3	Ensure life cycle independence	*	**
4	Secure HGW's baseline services entity	*	**
5	Open entities for 3 <sup>rd</sup> party and "unfriendly open source"	*	**
6	Allow different kernel versions to run in parallel		**
7	Allow different OS to run in parallel		**
8	Leverage Cloud-based on-demand HW resources (CPU, memory, storage)	*	*
	<b>Essential Characteristics</b>		
1	Memory Overhead	**	
2	Performance	**	
3	Boot time	**	*
4	No specific HW requirements	**	
5	Ease of implementation	*	
6	Ease of SW migration	*	**

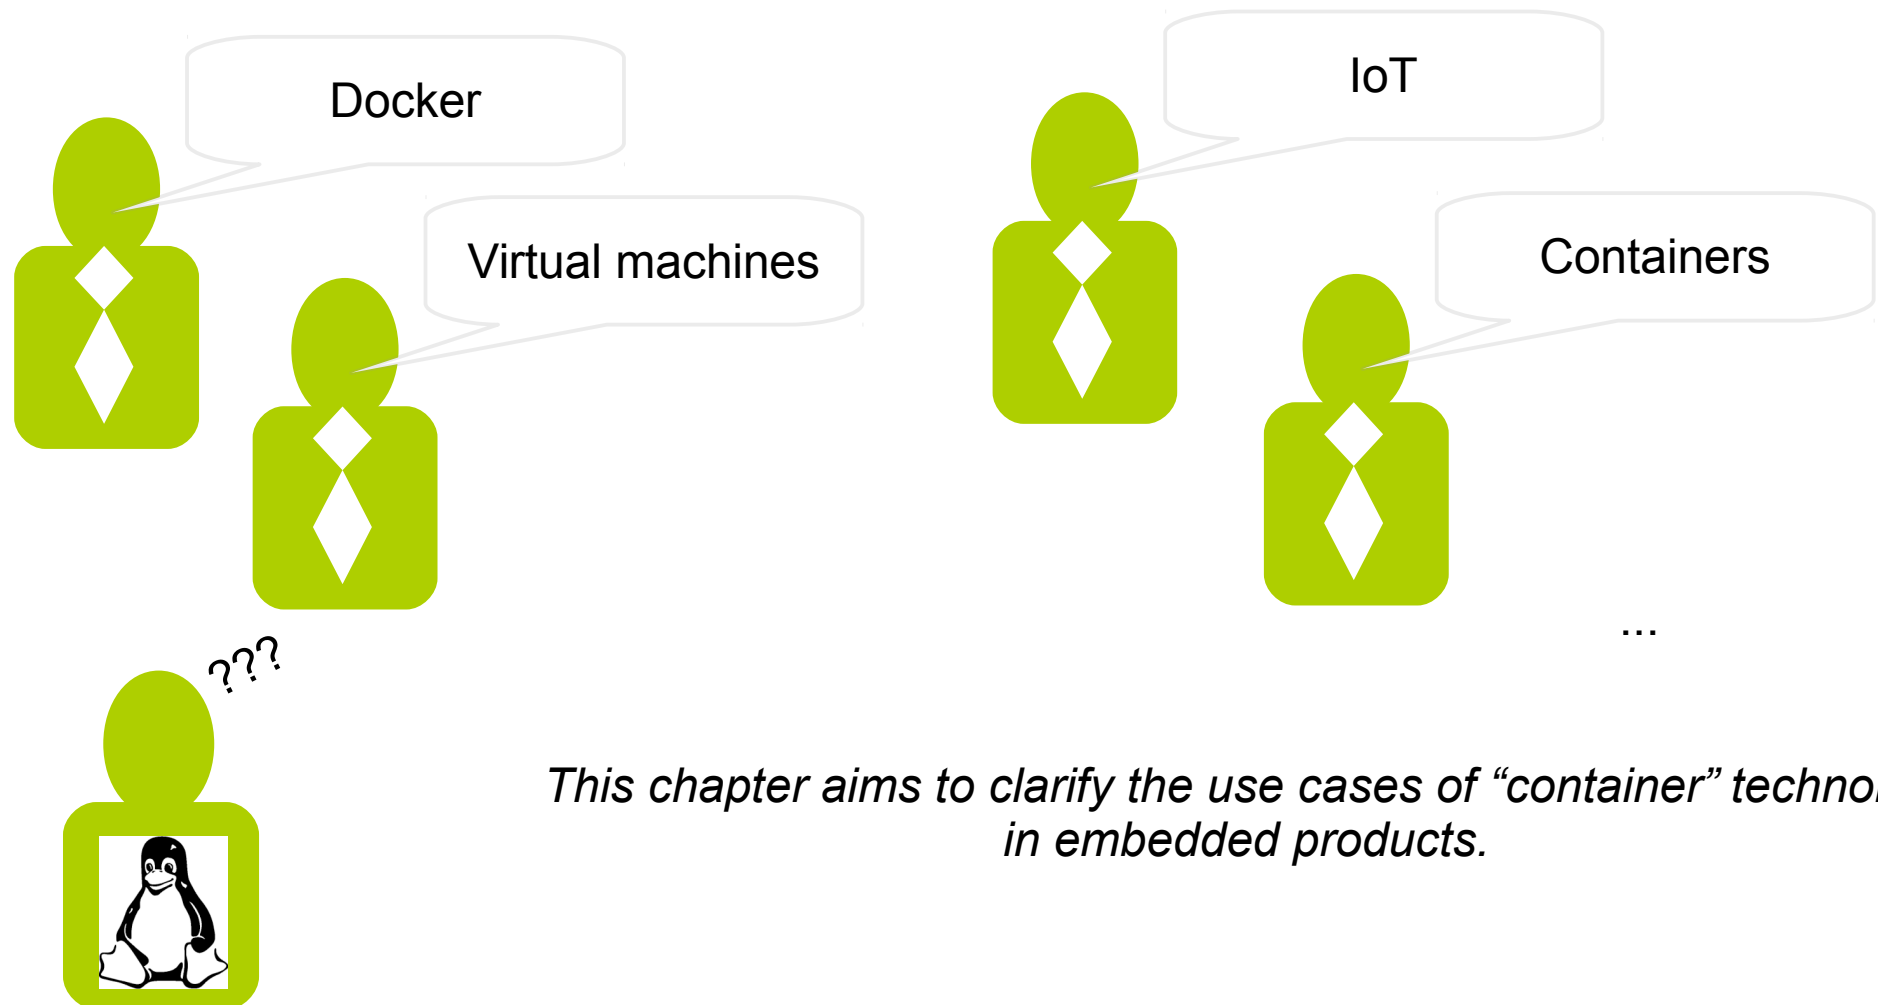
## Containers or HW Virtualisation? → **Both**

- Containers
  - address HGW challenges partially... but in a **very efficient** way
  - “lightweight virtualisation” - **available now**
- HW Virtualisation
  - addresses more challenges... but with additional HW/SW burden/cost
  - **not yet available**... but on some router SoC roadmaps
- Conclusions
  - Containers are an obvious step forward – available now
  - **Mix of containers and HW virtualisation** looks like a **future winning solution** addressing the challenges in the most efficient way
    - Example: an open application environment in a dedicated “open” Virtual Machine with each app packaged in its own container
  - To become reality, HW Virtualisation requires effort from chipset manufacturers

## **Part 3**

# **Containers on Embedded Platforms**

# Which 'hype' are we going to implement today?

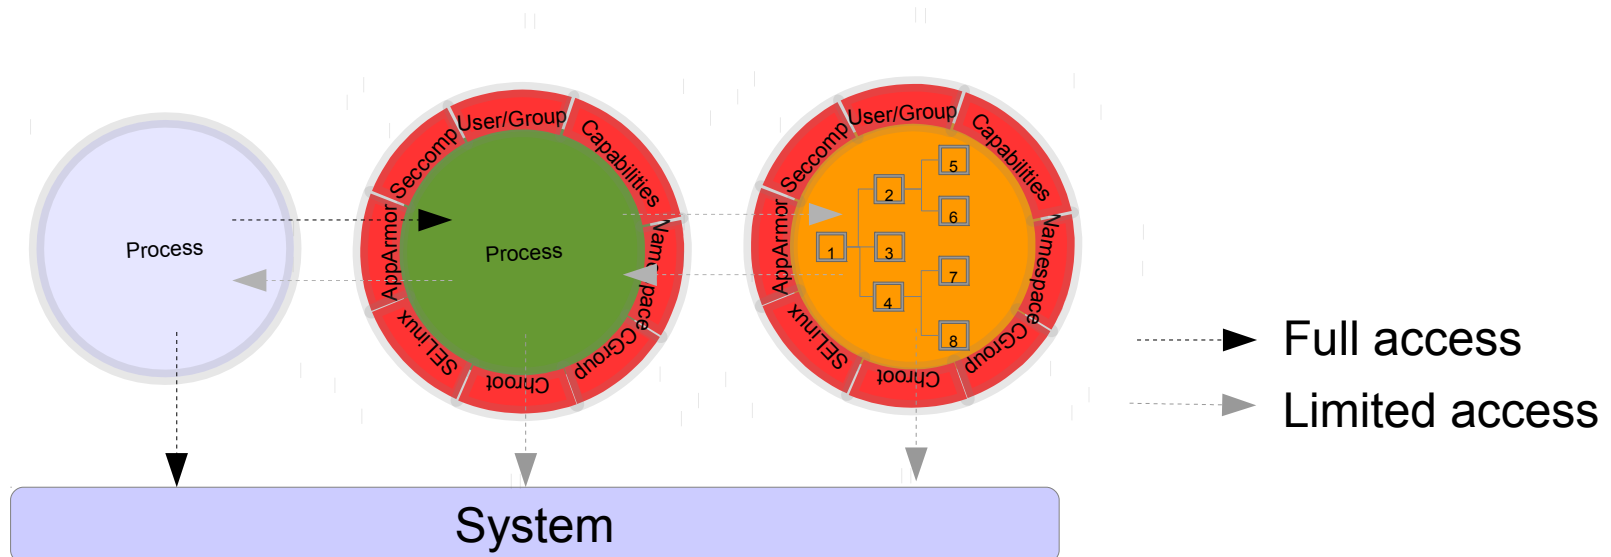


*This chapter aims to clarify the use cases of “container” technology in embedded products.*

# Implementing containers on embedded platforms

There are 2 possible **ways** of **implementing** container technology:

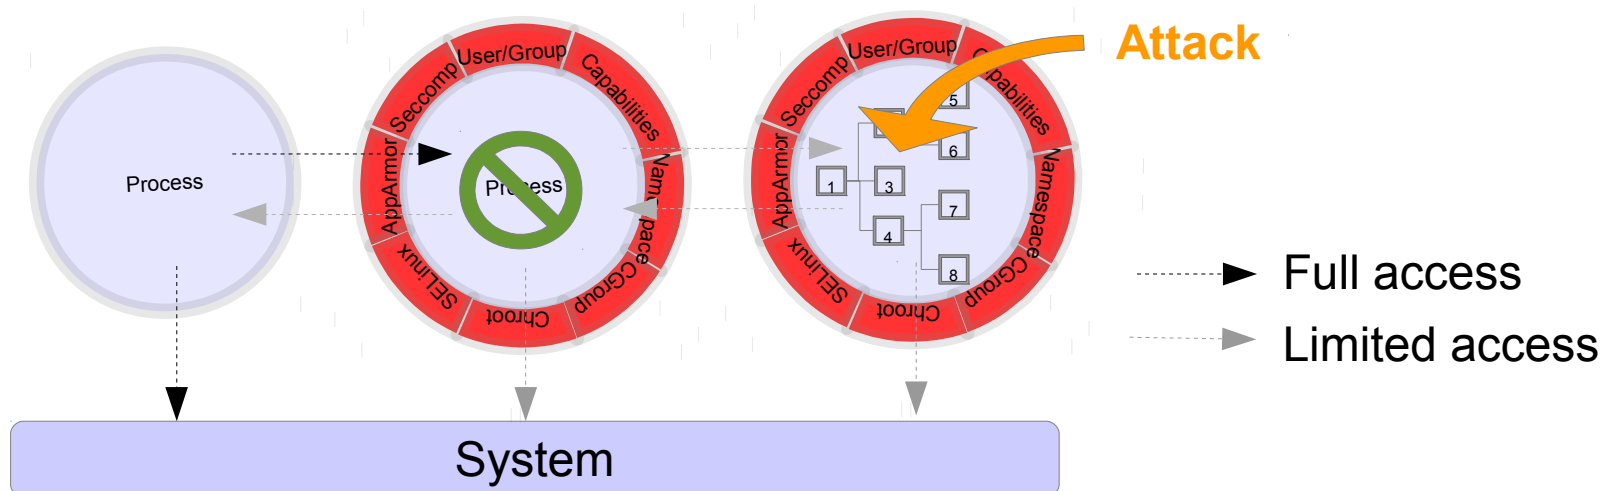
- **Single process** per container
- Containers as **virtual machines**



# Technical goals of container technology

These can be used to achieve 2 main **goals**:

- Increase **robustness** of a software stack
  - Running processes in a contained environment **shields the system** from bugs within those processes.
- Improve **security** of a software stack
  - Running a process in a contained environment isolates processes from the other system components, **reducing the attack surface**.

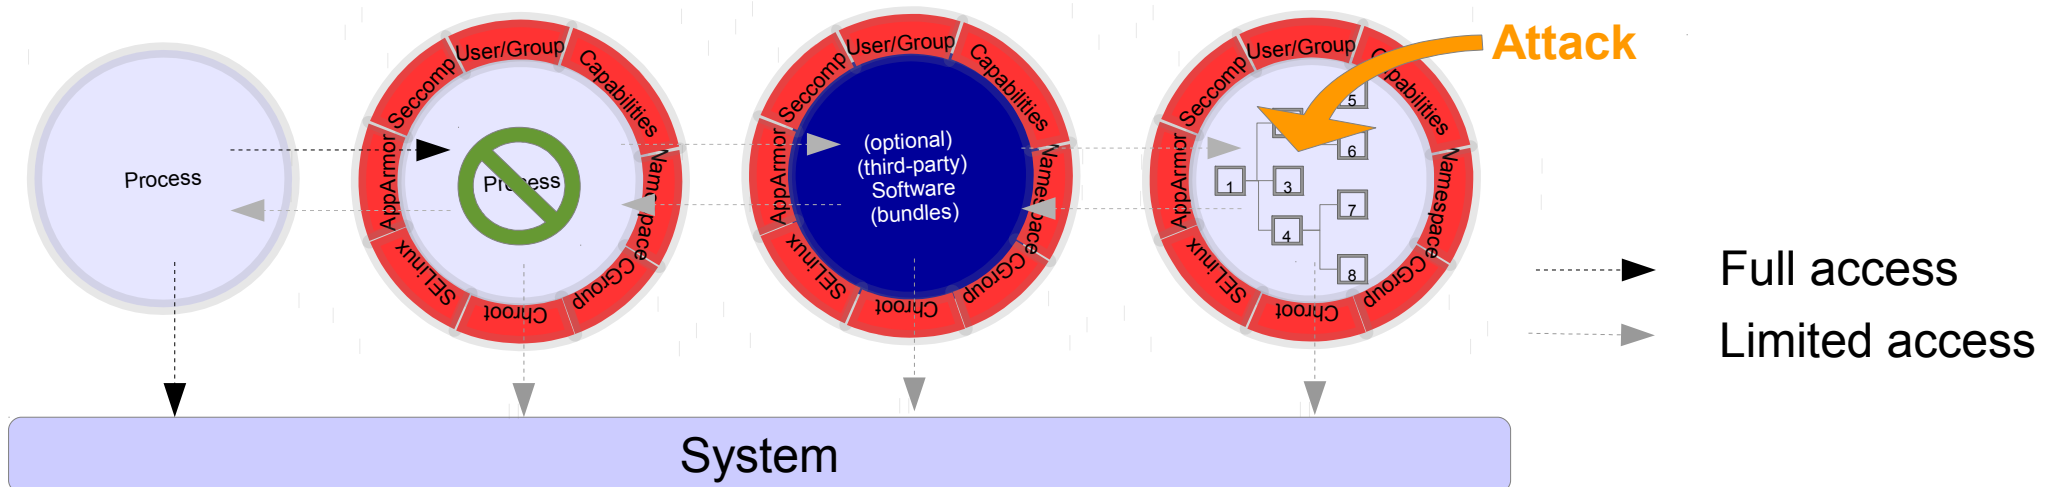




# Embedded use cases of containers

This means you can use containers for different **use cases**:

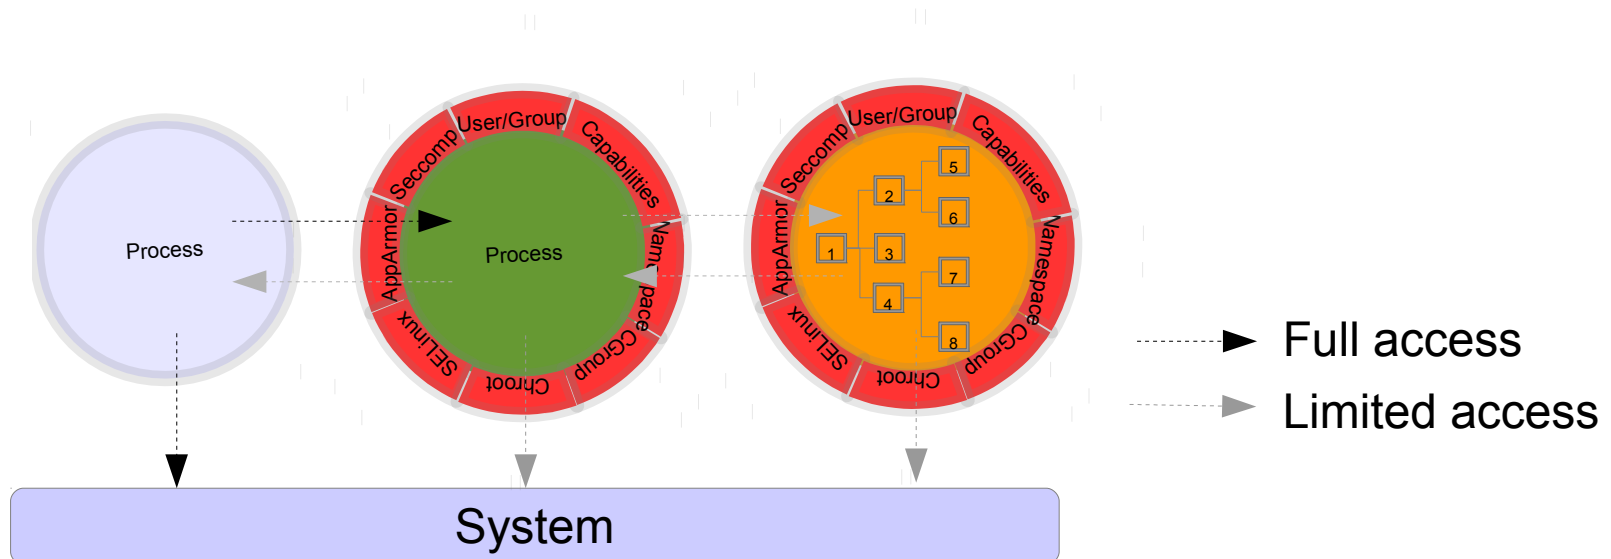
- Improve **robustness** of existing software stack.
  - e.g. Isolate specific processes that might impact system stability
- Improve **security** of existing software stack.
  - e.g. Secure processes that can be reached externally
- Integrate **(third-party) (optional) software (bundles)**.
  - Put binary third-party deliveries into sandbox environment
  - Enables app-store like concepts (of complex software stacks)



# Single process vs virtual machine approach

**Single process** per container approach:

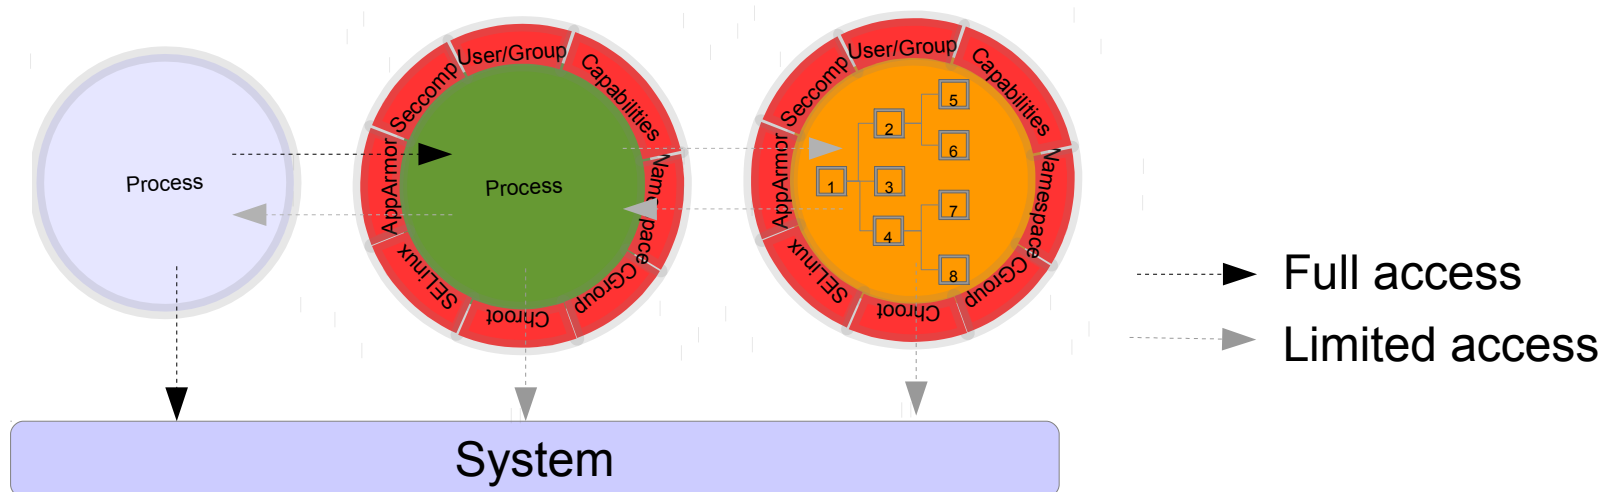
- + Configuration can be really tailored for that executable.
- + Optimal security and shielding
- Runtime dependencies are an issue
- No 'logical' packaging (think app-store)



# Single process vs virtual machine approach

**Virtual machines** per container approach:

- + Possibility to create 'logical' packages
- + Potentially resolves a lot of run-time dependencies
- Applications within container are not protected against each other
- Shielding depends on the weakest link
  - e.g. if one process requires privileges to mount a USB disk, all the other elements in the container have the same ability



## Side note

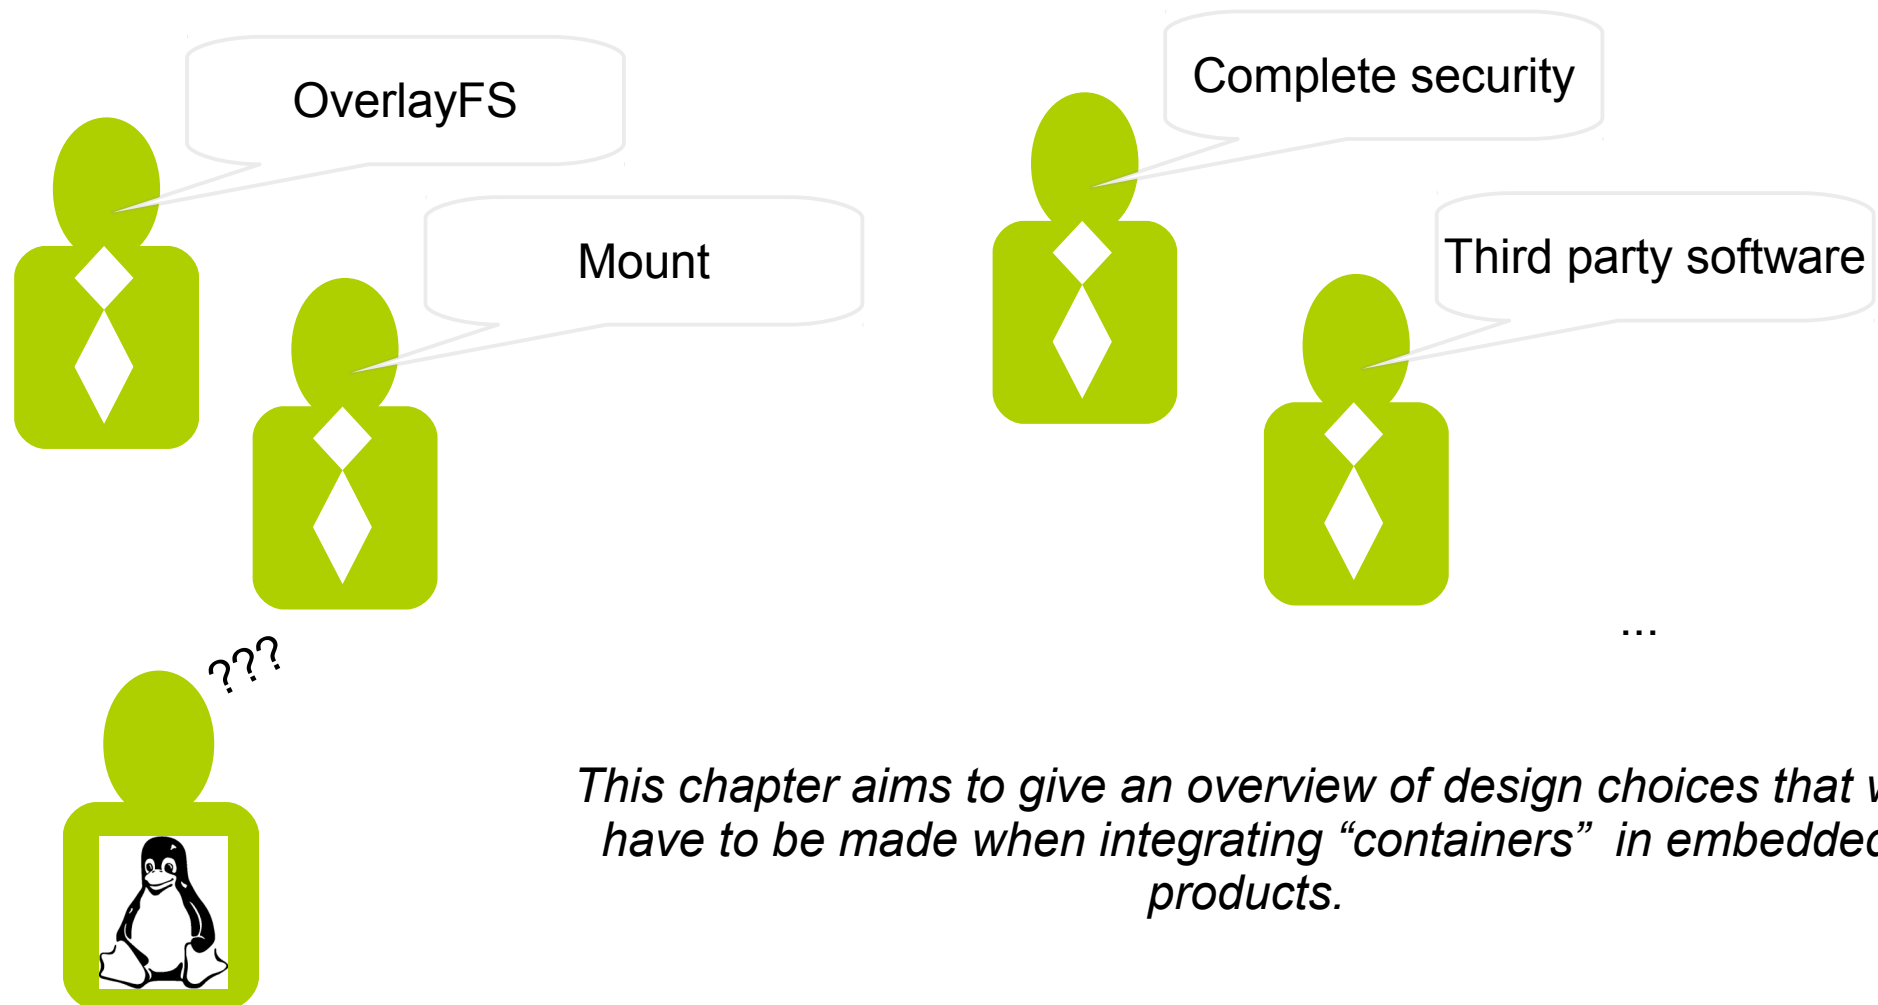
Why use container technology (Docker/LXC) and not use the kernel API's directly in the source of the different services ?

- ⇒ Adapting each service is time consuming and error prone.
- ⇒ Container frameworks offer a uniform interface to configure and manage a large variety of different settings.
- ⇒ The uniform interface makes it easier to integrate into an existing build system.

# **Part 4**

## **Container Design Choices**

# Just run it inside a container, right?



## Analyse your setup

- Which Linux kernel version?
- Which container facilities apply?
- What are the device limitations? (CPU/RAM/flash)

# Which Linux kernel version?

Kernel	Mount NS	UTS NS	IPC NS	PID NS	Network NS	User NS	CGroup v1	CGroup v2	APP Armor	SE Linux	Seccomp
2.4.19	X										
2.6.0	X									X	
2.6.12	X									X	X
2.6.19	X	X	X							X	X
2.6.24	X	X	X	X	X		X			X	X
2.6.36	X	X	X	X	X		X		X	X	X
3.8	X	X	X	X	X	X	X		X	X	X
4.5	X	X	X	X	X	X		X	X	X	X



## Which container facilities apply?

- The user/group policy
- Capabilities
- Namespaces: Mount, UTS, IPC, PID, Network, User
- CGroups
- Chroot
- SELinux/AppArmor
- Seccomp

Some configuration options are easy to configure:

- The user/group policy
- Capabilities
- Namespaces: Mount, UTS, IPC, PID, Network, User
- CGroups
- Chroot
- SELinux/AppArmor
- Seccomp

Others are more difficult to use and configure:

Network - File system - Users

- The user/group policy
- Capabilities
- Namespaces: Mount, UTS, IPC, PID, Network, User
- CGroups
- Chroot
- SELinux/AppArmor
- Seccomp

## Configuring your network: types


- empty Only loopback interface in container
- phys Moves the host interface to the container namespace
- vlan VLAN on top of interface
- none Shared network namespace between host and container
- veth Virtual Ethernet interface
- macvlan Multiple MAC/IP on the same physical interface
  - 3 modes : VEPA, private, bridge

## Configuring your network: empty/phys/vlan

- Empty: When container does not need networking
- Phys: Might be used for complex use case e.g. Ethernet over USB etc.
- VLAN: Use cases limited in consumer environment

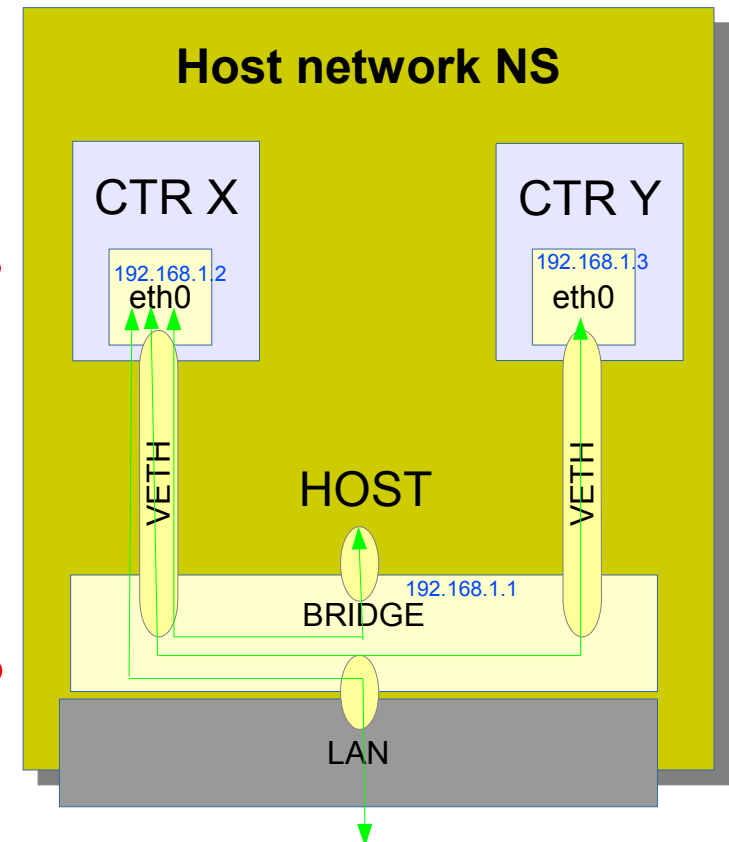
## Configuring your network: none

- + All host interfaces reachable in container
- + No performance penalty
- + Easy to listen on host IP address
- Less secure
  - \* container can access all interfaces
  - \* container can block ports
  - \* containers & firewall
  - \* ...
- Might be complex for third parties developing code in the container
- Network service configuration is shared; how to block port usage in container?



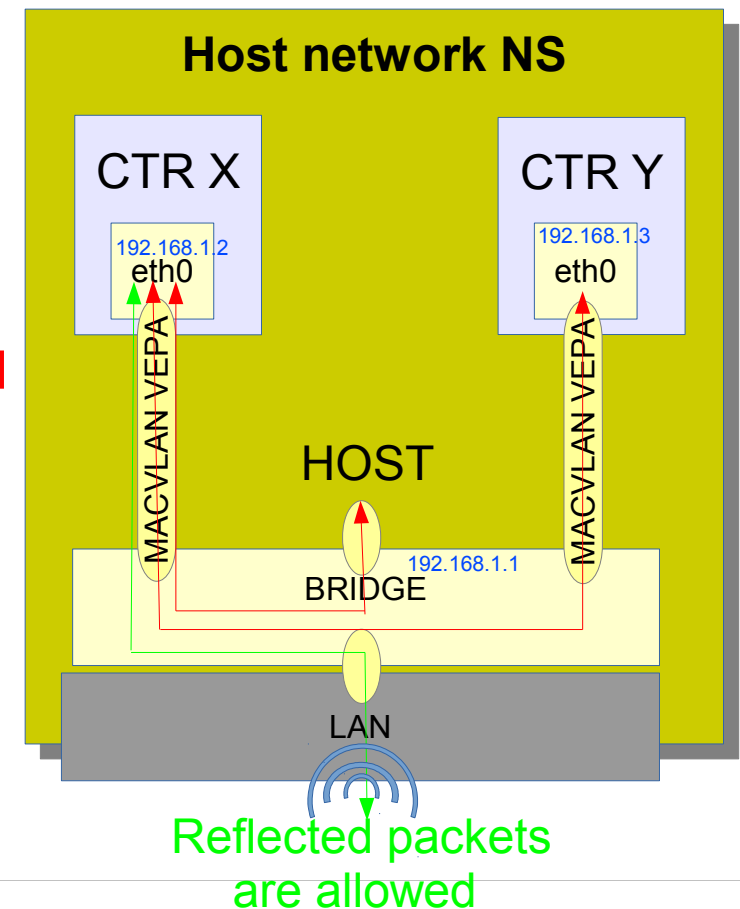
## Configuring your network: veth

- + Typically connected to software bridge on host
- + Container can access host + other containers
- + Containers can share a network namespace
- Security: container can access host
- Security: container can access other containers
- Security: firewall configuration needed
- Performance penalty caveats
- Port forwarding needed in case a service should be reachable on host IP
- Traffic from the container will have a different IP address.



# Configuring your network: Macvlan VEPA

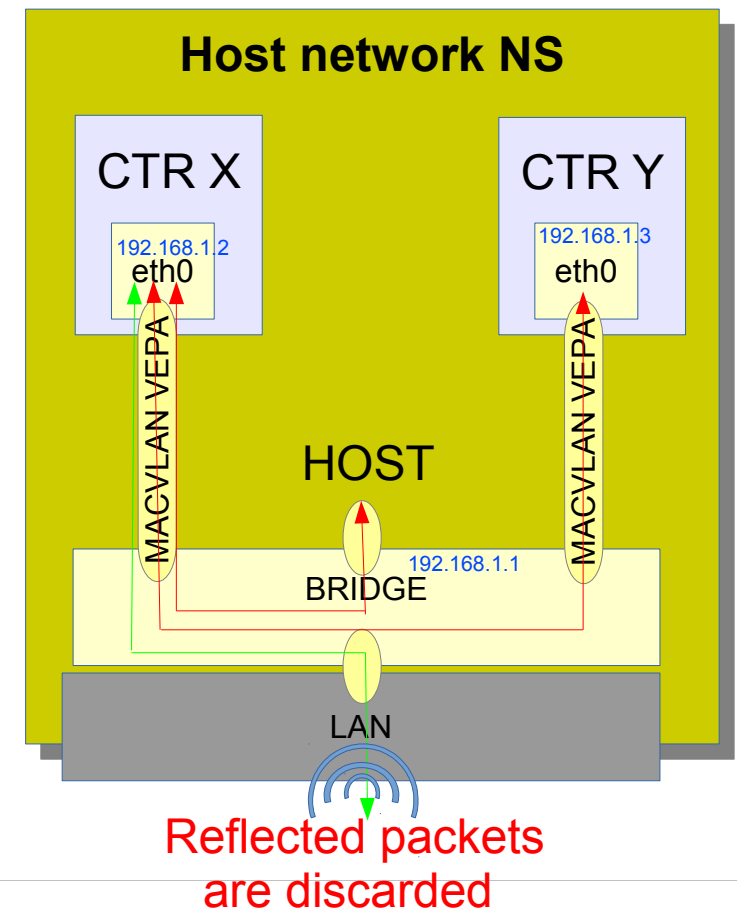
- + Packets directly written to interface, no performance penalty
- + Container can not access host + other containers unless packets are reflected or hairpin mode is used
- + Security : everything is blocked by default
- Only on (local) interfaces where you can add extra MAC's and IP's dynamically
- Port forwarding needed in case a service should be reachable on host IP
- Traffic from the container will have a different IP address.





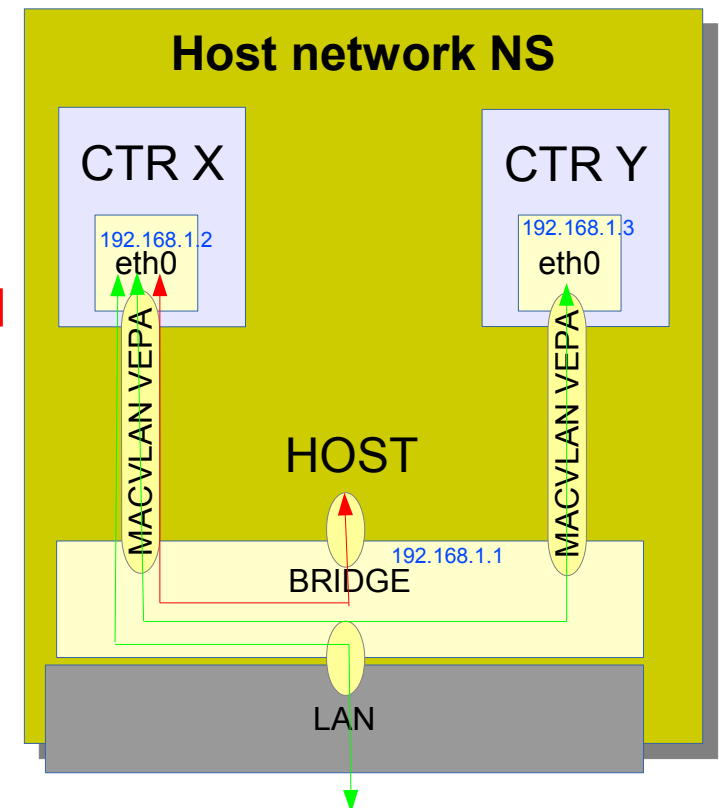
## Configuring your network: Macvlan private

Exactly the same as VEPA, but reflected packets are discarded



# Configuring your network: Macvlan bridge

- + Packets directly written to interface, no performance penalty
- + Container cannot access host unless hairpin mode is used
- Security: containers can talk to each other
- Only on (local) interfaces where you can add extra MAC's and IP's dynamically
- Port forwarding needed in case a service should be reachable on host IP
- Traffic from the container will have a different IP address.



## Defining your file system: RootFS requirements

- Different (often conflicting) requirements pop up:
  - As small as possible (flash + memory)
  - As performant as possible
  - Stand-alone capabilities
  - Security:
    - Preferably read-only / executable
    - Signature check at load
    - Correct file permissions
  - ...

# Defining your file system: RootFS technologies

- Container technologies:
  - File permissions
  - Chroot
  - SELinux/AppArmor
- Closely linked technologies:
  - Squashfs
  - Hardlinks
  - OverlayFS (or alternatives)

# Defining your file system: RootFS for **securing** an existing software stack

- Considerations for securing the rootfs
  - Rootfs is defined at build time of host
  - Hardlinks are possible, but watch out for file permissions!
  - Possible to **optimise** for **flash** size (no duplicate files)
  - Possible to **optimise** for **RAM** size (no duplicate files cached)
  - **Native performance** possible (no cache misses)
  - Fine tuning of SELinux/APPArmor possible (and really needed if you want to do it right)

## Defining your file system: RootFS for software bundles

- Considerations for creating the rootfs for software bundles
  - Rootfs is **not always** defined at build time of host
  - Hardlinks are possible, but more difficult to use!
  - **Stand alone capabilities** are more important, but have a cost on RAM/flash/performance.
  - Fine tuning of SELinux/APPArmor possible

## Defining your file system: RootFS implementation options

- Rootfs shared with host, shielded by SELinux/Apparmor
  - + For real security SELinux/APParmor mandatory anyway
  - + No performance/memory impact
  - File permissions in combination with user namespace?
  - Containers can see all files
  - Difficult to maintain

### CPE rootfs

Shielded RootFS 1

Shielded RootFS 2

# Defining your file system: RootFS implementation options

- Chroot to a hardlinked rootfs per container
  - + No performance/memory impact
  - + Cleaner view from container point of view
  - Hardlinks typically created at build time
  - Hardlink file permissions combination with user namespace?
  - For real security SELinux/APPArmor needed anyway

## CPE rootfs

hardlinks1	Standalone RootFS 1
hardlinks2	Standalone RootFS 2



# Defining your file system: RootFS implementation options

- Chroot to a duplicated rootfs per container
  - + No issues with file permissions in combination with user namespace
  - + Cleaner view from container point of view
  - + Minimal dependency on host rootfs
  - Performance/memory impact
  - Note: For real security SELinux/AppArmor needed anyway


## CPE rootfs

Standalone RootFS 1

Standalone RootFS 2

# Defining your file system: RootFS implementation options

- Share a base rootfs, use overlayFS as addon
  - + Solves memory issue of using separate rootfs partially
  - + Minimal dependency on host rootfs
  - + Standardisation of container rootfs possible
  - OverlayFS needs to be managed
  - Note: for real security SELinux/APPArmor needed anyway



# Defining your file system: RootFS implementation options overview

	Memory usage	Performance	Standalone	Manageability
Stand alone package	--	-	++	+
Shared base package	+	-	+	+
Shared with host – hard links	++	++	--	-

→ Which approach to choose depends on the goal you want to achieve

# Defining your file system: Tmpfs

- Requirements:
  - Read-write - preferably no execute
  - Limited in size
- Possible solutions:
  - Tmpfs mount in container configuration
    - ⇒ Easier to manage
  - Tmpfs mount in container firmware
    - ⇒ More freedom for container developers, tmpfs is counted into cgroup limitations

# Defining your file system: Persistent storage

- Requirements:
  - Read-write - no execute
  - Limited in size
- Possible solutions:
  - Bind mount R/W directory from host
  - Bind mount R/W directory from file (containing a FS) from host
  - Mount native partition in flash
  - Provide key/value store service on host
  - OverlayFS

# Defining your file system: Persistent storage

- Bind mount R/W directory from host:
  - + No runtime overhead
  - + Used by Docker, which allows you to 'name' a volume and share over different containers
  - Migration of data is difficult
  - No limitation possible on size:
    - \* user quota works in Linux per owner/group, not per directory
    - \* <https://code.google.com/p/fusequota/> but this is GPLv3
- Bind mount R/W directory from file (containing a FS) from host:
  - + Easy to limit in size
  - Possibly stability issues (how to control which data was really written to the flash)
  - Migration of data is difficult
  - Runtime overhead

# Defining your file system: Persistent storage

- Mount native partition in flash:
  - + No runtime overhead
  - + Easy to limit in size
  - Dynamic management of flash partitions needed
  - Migration of data is difficult
- Provide key/value store service on host :
  - + Easy to limit in size
  - + Migration of data can be easily managed in centralised way
  - + Can be used to push default settings to container as well
  - Runtime overhead
  - Complicated
  - Not easy to use
- Overlayfs (can be used in combination with the previous techniques)

## Defining the user namespace

- Map virtual user ID's with root-like privileges to real user ID's on the host (« unprivileged containers »)  
map real user 100000 → 165536 to virtual user 0 → 65536
- possible to be 'root' in container without being root on host
- root user in container gains extra capabilities within that container



# **Part 5**

## **Case Studies**

## Case 1: Integration of a NAS software stack

- Main use cases:
  - isolation from the router stack, mostly CPU, RAM, flash, I/O
  - security is not the main goal; NAS stack is semi-trusted
  - independent upgrade cycle from the main router firmware
- « virtual machine » model
  - the container has its own rootfs
  - a list of processes is started at container init

## Case 1: Integration of a NAS software stack: Implementation choices

- Single container, always present.
- (Latest version of the) rootfs is embedded in the software image of the router, as a subdirectory of the router's rootfs.
- Upgrade of the container via TR-069 or web UI is possible. The upgrade file is:
  - stored in a flash partition
  - digitally signed, and checked at installation and at every boot
  - contains a squashfs rootfs, mounted through loopback
- Two network interfaces in the NAS « vm »:
  - veth in a private bridge with the host, firewalled, for communication with the host and (NAT'ed) the WAN
  - veth in the LAN bridge, with its own MAC and LAN IP address
- USB storage devices are managed by the NAS « vm » with udev rules
- tmpfs for temporary storage, dedicated flash partition for persistent storage

## Case 2: Untrusted third party application store

- Main use cases:
  - third party application deployment on the router
  - isolation from the router stack at every possible level
  - security is critical; the application is untrusted
- Lightweight « virtual machine » model
  - the container's rootfs overlays an (upgradable) « base » filesystem with a set of standard libraries and utilities
  - one or more processes are started at container init

## Case 2: Untrusted third party application store: Implementation choices

- any number of containers can be installed
- unprivileged containers are mandatory
- the container image file:
  - is stored in a flash partition
  - is digitally signed, and checked at installation and at every boot
  - contains a squashfs rootfs, mounted through loopback
  - contains a large list of customised configuration settings:
    - network configuration, RAM, CPU, I/O, firewall settings
    - fine-grained access control on management IPC mechanism

... perhaps this use case is better addressed with HW virtualisation.



---

For more information about  
**SoftAtHome**

**Corporate Headquarter**

9 rue du Débarcadère  
92700 Colombes  
France

**United Arab Emirates office**

Building 6EA office 207  
P.O. Box 371438,  
Dubai Airport Freezone  
United Arab Emirates

**Belgium office**

Vaartdijk 3 701  
3018 Wijnmaal  
Belgium